



Trip+: Benchmarking Agents in Personalized Interactive Travel Planning




Junle Chen^{1,4}, Wei Chen^{1,2*}, Yehong Xu^{1,4}, Zhengjun Huang^{1,4}, Yuqian Wu³,
Zhoujin Tian^{1,4}, Kai Wang², Rui Chen², Xiaofang Zhou^{1,4}

¹ The Hong Kong University of Science and Technology ² Tencent Hy

³ The Hong Kong University of Science and Technology (Guangzhou) ⁴ Guangzhou HKUST Fok Ying Tung Research Institute

Motivation & contributions



Why both matter

Personalization: travelers differ in interests, budgets and physical limits.
Interaction: needs evolve, and disruptions require plan revisions.

The gap


Existing benchmarks underexplore **rich personalization** combined with **diverse, long-horizon interactions**.

Trip+ links both

Profile-aware plan revision, stateful checks and traveler simulation.

Interaction reliability

Gemini-3.1-Pro-Preview (all four scenarios)




Takeaway


Better interaction requires revising plans reliably as constraints grow.

Benchmark design & evaluation

153 tasks, 570 turns, 11 profiles and 40 cities.
Agents use 11 travel tools to plan and revise trips.




Stateful evaluation separates fulfillment of new requests from preservation of earlier requirements. Four LLM judges replay final itineraries to assess traveler experience.



Traveler experience

Gemini-3.1-Pro-Preview

Plan quality (0-1)



Pacing and experience burden



Takeaway

Better personalization depends on pacing and burden control.

Main results

18 models with the same agent scaffold

Best observed Plan Avg.: 0.7331 (Gemini-3.1-Pro-Preview). Best observed User Sim.: 0.5184 (GPT-5.4).

Model	End-to-End Interaction			Valid Plan Quality			Summary		
	Resp. Acc.	Req. Fulfill	Intent Pres.	Feas.	Hard Cons.	Soft Pref.	User Sim.	Plan Avg.	Win(%)
Frontier models									
Gemini-3.1-Pro-Preview	0.8544	0.7382	0.7697	0.8812	0.9047	0.6389	0.5077	0.7331	42.9
Gemini-3-Flash-Preview	0.8667	0.6885	0.6532	0.7099	0.7179	0.6172	0.4862	0.6328	0.0
GPT-5.4	0.7544	0.5585	0.5489	0.7669	0.8110	0.5854	0.5184	0.6704	14.3
GPT-5.4-Mini	0.7807	0.4879	0.4221	0.6455	0.6126	0.6207	0.5085	0.5968	0.0
GLM-5.1	0.8702	0.7214	0.7301	0.8463	0.7973	0.6136	0.5014	0.6896	0.0
Kimi-K2.6	0.8439	0.6727	0.6811	0.7472	0.7984	0.6380	0.5164	0.6750	0.0
Doubao-Seed-2.0-Pro	0.8246	0.6811	0.6814	0.6731	0.8006	0.6031	0.5128	0.6474	0.0
DeepSeek-V4-Pro	0.7912	0.6657	0.7139	0.6755	0.8375	0.6258	0.5100	0.6622	0.0
DeepSeek-V3.2	0.7702	0.5418	0.5566	0.6461	0.7162	0.5842	0.4755	0.6055	0.0
Hy3-preview	0.8930	0.6708	0.7640	0.7223	0.6476	0.5957	0.4866	0.6130	0.0
MiniMax-M2.7	0.7772	0.5766	0.4921	0.5596	0.6326	0.6076	0.4992	0.5748	0.0
Lightweight models									
Gemma-4-31B	0.9035	0.7465	0.7490	0.7427	0.7632	0.5731	0.4923	0.6428	28.6
Gemma-4-26B-A4B	0.8386	0.5989	0.5594	0.6761	0.6925	0.5940	0.4916	0.6136	0.0
Qwen3.5-27B	0.8491	0.6309	0.6218	0.7393	0.7256	0.6375	0.4821	0.6461	0.0
Qwen3.5-122B-A10B-FP8	0.3368	0.2033	0.1028	0.6935	0.7235	0.6319	0.5050	0.6385	0.0
Qwen3.6-27B	0.5649	0.4350	0.3756	0.7521	0.7225	0.6626	0.5167	0.6635	14.3
Qwen3.6-35B-A3B	0.6912	0.4554	0.2463	0.4292	0.4588	0.5839	0.5141	0.4965	0.0
GLM-4-32B	0.5053	0.4656	0.3155	0.5236	0.5605	0.6141	0.4791	0.5443	0.0

Takeaway: Feasible itineraries are not necessarily user-aligned.

Reliable revision and traveler-aware pacing remain open challenges

Scope: synthetic English dialogues in 40 Chinese cities. Simulated comfort is a proxy for traveler experience.

Website: <https://junle-chen.github.io/trip-plus-site/>

Code & data: github.com/junle-chen/trip-plus

Email: jchenkg@connect.ust.hk



Code & data