



# Trip+: Benchmarking Agents in Personalized Interactive Travel Planning




Junle Chen<sup>1,4</sup>, Wei Chen<sup>1,2\*</sup>, Yehong Xu<sup>1,4</sup>, Zhengjun Huang<sup>1,4</sup>, Yuqian Wu<sup>3</sup>,  
Zhoujin Tian<sup>1,4</sup>, Kai Wang<sup>2</sup>, Rui Chen<sup>2</sup>, Xiaofang Zhou<sup>1,4</sup>

<sup>1</sup> The Hong Kong University of Science and Technology <sup>2</sup> Tencent Hy

<sup>3</sup> The Hong Kong University of Science and Technology (Guangzhou) <sup>4</sup> Guangzhou HKUST Fok Ying Tung Research Institute

## Motivation & contributions



### Why both matter

**Personalization:** preferences differ across travelers.  
**Interaction:** needs evolve; disruptions require repeated plan revisions.

### The gap


Existing benchmarks underexplore **rich personalization** across diverse interactions. Agents must meet new requests while respecting traveler preferences and **earlier requirements**.

### Trip+ evaluates both

Rich profiles, diverse interactions, stateful intent checks and traveler simulation.

## Interaction reliability

Gemini-3.1-Pro-Preview (all four scenarios)





### Takeaway

**Better interaction requires revising plans reliably as constraints grow.**

## Benchmark design & evaluation

**153 tasks, 570 turns, 11 profiles and 40 cities.**  
Agents use 11 travel tools to plan and revise trips.




Stateful evaluation separates fulfillment of new requests from preservation of earlier requirements. Four LLM judges replay final itineraries to assess traveler experience.


## Traveler experience

Gemini-3.1-Pro-Preview

Plan quality (0-1)



### Pacing and experience burden



### Takeaway

**Better personalization depends on pacing and burden control.**

## Main results

18 models with the same agent scaffold

Best observed Plan Avg.: 0.7331 (Gemini-3.1-Pro-Preview). Best observed User Sim.: 0.5184 (GPT-5.4).

| Model                     | End-to-End Interaction |              |              | Valid Plan Quality |            |            | Summary   |           |        |
|---------------------------|------------------------|--------------|--------------|--------------------|------------|------------|-----------|-----------|--------|
|                           | Resp. Acc.             | Req. Fulfill | Intent Pres. | Feas.              | Hard Cons. | Soft Pref. | User Sim. | Plan Avg. | Win(%) |
| <b>Frontier models</b>    |                        |              |              |                    |            |            |           |           |        |
| Gemini-3.1-Pro-Preview    | 0.8544                 | 0.7382       | 0.7697       | 0.8812             | 0.9047     | 0.6389     | 0.5077    | 0.7331    | 42.9   |
| Gemini-3-Flash-Preview    | 0.8667                 | 0.6885       | 0.6532       | 0.7099             | 0.7179     | 0.6172     | 0.4862    | 0.6328    | 0.0    |
| GPT-5.4                   | 0.7544                 | 0.5585       | 0.5489       | 0.7669             | 0.8110     | 0.5854     | 0.5184    | 0.6704    | 14.3   |
| GPT-5.4-Mini              | 0.7807                 | 0.4879       | 0.4221       | 0.6455             | 0.6126     | 0.6207     | 0.5085    | 0.5968    | 0.0    |
| GLM-5.1                   | 0.8702                 | 0.7214       | 0.7301       | 0.8463             | 0.7973     | 0.6136     | 0.5014    | 0.6896    | 0.0    |
| Kimi-K2.6                 | 0.8439                 | 0.6727       | 0.6811       | 0.7472             | 0.7984     | 0.6380     | 0.5164    | 0.6750    | 0.0    |
| Doubao-Seed-2.0-Pro       | 0.8246                 | 0.6811       | 0.6814       | 0.6731             | 0.8006     | 0.6031     | 0.5128    | 0.6474    | 0.0    |
| DeepSeek-V4-Pro           | 0.7912                 | 0.6657       | 0.7139       | 0.6755             | 0.8375     | 0.6258     | 0.5100    | 0.6622    | 0.0    |
| DeepSeek-V3.2             | 0.7702                 | 0.5418       | 0.5566       | 0.6461             | 0.7162     | 0.5842     | 0.4755    | 0.6055    | 0.0    |
| Hy3-preview               | 0.8930                 | 0.6708       | 0.7640       | 0.7223             | 0.6476     | 0.5957     | 0.4866    | 0.6130    | 0.0    |
| MiniMax-M2.7              | 0.7772                 | 0.5766       | 0.4921       | 0.5596             | 0.6326     | 0.6076     | 0.4992    | 0.5748    | 0.0    |
| <b>Lightweight models</b> |                        |              |              |                    |            |            |           |           |        |
| Gemma-4-31B               | 0.9035                 | 0.7465       | 0.7490       | 0.7427             | 0.7632     | 0.5731     | 0.4923    | 0.6428    | 28.6   |
| Gemma-4-26B-A4B           | 0.8386                 | 0.5989       | 0.5594       | 0.6761             | 0.6925     | 0.5940     | 0.4916    | 0.6136    | 0.0    |
| Qwen3.5-27B               | 0.8491                 | 0.6309       | 0.6218       | 0.7393             | 0.7256     | 0.6375     | 0.4821    | 0.6461    | 0.0    |
| Qwen3.5-122B-A10B-FP8     | 0.3368                 | 0.2033       | 0.1028       | 0.6935             | 0.7235     | 0.6319     | 0.5050    | 0.6385    | 0.0    |
| Qwen3.6-27B               | 0.5649                 | 0.4350       | 0.3756       | 0.7521             | 0.7225     | 0.6626     | 0.5167    | 0.6635    | 14.3   |
| Qwen3.6-35B-A3B           | 0.6912                 | 0.4554       | 0.2463       | 0.4292             | 0.4588     | 0.5839     | 0.5141    | 0.4965    | 0.0    |
| GLM-4-32B                 | 0.5053                 | 0.4656       | 0.3155       | 0.5236             | 0.5605     | 0.6141     | 0.4791    | 0.5443    | 0.0    |

**Takeaway: Feasible itineraries are not necessarily user-aligned.**

## Reliable revision and traveler-aware pacing remain open challenges

Website: <https://junle-chen.github.io/trip-plus-site/>

Code & data: [github.com/junle-chen/trip-plus](https://github.com/junle-chen/trip-plus)

Email: [jchenkg@connect.ust.hk](mailto:jchenkg@connect.ust.hk)



Code & data